

Event-TimeRAF: Event-Aware Retrieval-Augmented Foundation Model for Explainable Air Quality Forecasting Under Concept Drift

Author Name, *Student Member, IEEE*

Abstract—Urban air-quality forecasting is important for public-health warnings, environmental management, and city-scale decision support. Time-series foundation models have recently shown promising zero-shot forecasting capability, and TimeRAF demonstrates that retrieval augmentation can supplement a frozen forecasting backbone with useful historical knowledge. However, PM_{2.5} variation is strongly affected by local context such as meteorology, holidays, high-wind events, heavy rain, wildfire activity, and other environmental disturbances. This paper introduces Event-TimeRAF, an event-aware retrieval-augmented framework for Los Angeles County PM_{2.5} forecasting. Event-TimeRAF builds a source-preserving knowledge base from official air-quality, weather, calendar, and storm-event records; retrieves causally eligible historical windows; fuses retrieved evidence with tabular forecasters and a frozen Chronos-Bolt foundation model; and generates evidence-grounded explanations under an explicit drift audit. We evaluate the framework on EPA AirData PM_{2.5}, NOAA Global Hourly weather, and hash-verified NOAA Storm Events records for 2019–2024. The strongest overall model is an XGBoost weather-calendar context baseline with MSE 26.185, while the full Event-TimeRAF XGBoost variant obtains MSE 26.712. The frozen Chronos-Bolt baseline obtains MSE 28.941, and retrieval-augmented Chronos reduces it to MSE 28.709. These results show that retrieval is beneficial in selected settings, especially for the frozen foundation-model fusion, but is not uniformly superior to strong supervised context features. Because NOAA Storm Events lacks machine-readable publication timestamps, event-aware results are reported as retrospective event-start sensitivity results rather than strict real-time event forecasts.

Index Terms—Air quality forecasting, PM_{2.5}, retrieval-augmented forecasting, time-series foundation models, TimeRAF, concept drift, explainable forecasting.

I. INTRODUCTION

Fine particulate matter (PM_{2.5}) is a major urban air-quality risk because it is associated with respiratory and cardiovascular health effects [1]. Accurate short-term PM_{2.5} forecasts can support public-health alerts, exposure reduction, transport planning, and environmental operations. In Los Angeles County, however, PM_{2.5} dynamics are difficult to model with pollutant history alone. Concentrations can change with wind speed, boundary-layer behavior, rainfall, seasonal activity, holidays, wildfire smoke, and unusual storm or high-wind events. These factors make the forecasting problem both temporal and contextual.

Recent time-series foundation models (TSFMs) aim to improve forecasting generalization by pre-training on large time-

series corpora and transferring the learned forecasting ability to new domains. Representative models include TimesFM, Moirai, MOMENT, Chronos, and Tiny Time Mixers [2]–[6]. In parallel, retrieval-augmented generation has shown that non-parametric memory can help foundation models access relevant external knowledge at inference time [7], [8]. TimeRAF combines these two ideas for zero-shot time-series forecasting by retrieving relevant historical sequences from a curated knowledge base and integrating them into a frozen TSFM through Channel Prompting [9].

Air-quality forecasting raises a domain-specific retrieval question that is not fully addressed by numerical time-series retrieval alone. A historical PM_{2.5} window may look similar to the current window but occur under different weather or event conditions. Conversely, a window with moderate numerical similarity may be more informative if it shares high-wind, rainfall, holiday, smoke, or stagnation context. This motivates a retrieval framework that treats source records, weather, calendar information, and distribution-shift indicators as part of the forecasting evidence rather than as unrelated metadata.

This paper introduces Event-TimeRAF, an event-aware retrieval-augmented framework for explainable PM_{2.5} forecasting under concept drift. The design is inspired by TimeRAF, but it is adapted to an official-source environmental forecasting setting. The implemented system uses leakage-safe retrieval over historical PM_{2.5} windows, weather and calendar context, source-recorded NOAA Storm Events, transparent drift indicators, direct XGBoost forecasters, a frozen Chronos-Bolt publication gate, and evidence-grounded explanations. The framework answers three practical questions: (1) how can official environmental records be converted into a retrieval-compatible knowledge base; (2) how should retrieval be constrained so that future target information is not leaked; and (3) when does retrieval help beyond strong weather-calendar supervised baselines?

The main contributions are as follows:

- We formulate an event-aware retrieval-augmented PM_{2.5} forecasting framework that combines time-series windows, weather summaries, calendar context, storm-event records, and drift indicators.
- We implement a Kaggle-compatible official-source pipeline using EPA AirData PM_{2.5}, NOAA Global Hourly weather, and hash-verified NOAA Storm Events records for Los Angeles County during 2019–2024.
- We define leakage-safe retrieval eligibility rules that prevent candidate windows from overlapping the query

context or future target.

- We evaluate random retrieval, cosine retrieval, event-aware hybrid retrieval, drift-aware feature fusion, frozen Chronos-Bolt, and retrieval-augmented Chronos-Bolt under one manifest-backed final run.
- We report a conservative empirical finding: retrieval improves the frozen Chronos-Bolt MSE slightly, but the best overall model is a strong XGBoost weather-calendar context baseline.

The final result is therefore not a uniformly positive retrieval story. Instead, it is an audited study showing where event-aware retrieval helps, where it does not, and what is required before such a method can be described as a strict operational event-aware air-quality forecaster.

II. LITERATURE REVIEW AND RELATED WORK

A. Time-Series Forecasting

Time-series forecasting has traditionally been studied through statistical models, machine-learning regressors, and deep neural networks. Autoregressive and seasonal models remain useful baselines because they are interpretable and efficient. Gradient-boosted tree models such as XGBoost and LightGBM are widely used for tabular forecasting with lagged and rolling features [10], [11]. Recurrent models such as LSTM and GRU can capture sequential dependence [12], [13]. Transformer-based models then expanded the long-sequence forecasting literature. Informer reduces attention cost for long inputs [14], Autoformer and FEDformer combine decomposition with long-range dependency modeling [15], [16], Crossformer and iTransformer focus on multivariate channel structure [17], [18], and TimesNet models temporal variation through 2D transformations [19]. At the same time, DLinear showed that simple linear decomposition models can outperform more complex Transformer variants on some long-term forecasting benchmarks [20]. PatchTST later used patching and channel-independent designs to improve long-lookback efficiency [21]. For air-quality forecasting, these approaches are helpful, but they may still be sensitive to distribution shifts and unobserved external events.

B. Time-Series Foundation Models

TSFMs aim to provide transferable forecasting ability across datasets and domains. Recent surveys describe TSFMs as a shift from task-specific forecasting toward reusable representations and zero/few-shot transfer [22], [23]. TimesFM uses a decoder-only architecture for time-series forecasting [2]. Moirai proposes a universal forecasting Transformer trained on a large-scale open time-series archive [3]. MOMENT provides a family of open time-series foundation models for general-purpose time-series tasks [4]. Chronos converts time-series values into tokens and trains language-model architectures for probabilistic forecasting [5]. TTM proposes compact pre-trained models for efficient zero/few-shot forecasting [6]. Lag-Llama is another probabilistic foundation model that uses a decoder-only architecture with lagged covariates and reports zero-shot transfer across unseen time-series domains [24]. These models motivate the use of pre-trained forecasting

backbones, but practical domain adaptation is still difficult when the downstream series responds strongly to local events and environmental conditions.

C. Retrieval-Augmented Forecasting

Retrieval-augmented generation first became popular in NLP by combining parametric models with non-parametric memory [7]. Dense Passage Retrieval (DPR) introduced dense dual-encoder retrieval for open-domain question answering [8]. TimeRAF adapts the retrieval-augmentation idea to zero-shot time-series forecasting [9]. It constructs a time-series knowledge base, retrieves top- k candidate windows using a learnable retriever, and combines them through Channel Prompting. The base paper reports TTM-Base as the main backbone, context length 512, forecast horizon 96 in most experiments, and $k = 8$ as the default number of retrieved candidates. Its ablations compare random retrieval, cosine retrieval, learnable retrieval, alternative integration mechanisms, candidate counts, knowledge-base choices, and backbone choices.

Event-TimeRAF follows the TimeRAF principle but changes both the knowledge source and the target domain. Instead of relying only on numerical time-series retrieval, it uses a hybrid knowledge base containing historical PM2.5 windows, weather summaries, calendar context, and event evidence. The implemented system is TimeRAF-inspired rather than a full reproduction: it evaluates random retrieval, cosine retrieval, event-aware hybrid retrieval, drift-aware feature fusion, frozen Chronos-Bolt, and retrieval-augmented Chronos-Bolt forecast fusion, while a learnable dual-encoder retriever and full Channel-Prompting neural integration remain future extensions.

D. Air-Quality Forecasting and Data Sources

PM2.5 forecasting often uses historical pollutant measurements together with meteorological factors such as temperature, humidity, wind speed, precipitation, and pressure. Recent PM2.5 studies and reviews show a shift from single-station recurrent models toward hybrid spatiotemporal models, attention mechanisms, graph neural networks, and interpretable workflows [25]–[28]. These studies motivate the use of meteorology, temporal context, and spatial/event structure, but they usually evaluate learned forecasters rather than source-audited retrieval evidence. The US Environmental Protection Agency publishes quality-controlled hourly observations through Air-Data and the Air Quality System [29]. NOAA’s Integrated Surface Database provides hourly meteorological observations from surface stations [30]. NOAA also publishes structured Storm Events records and historical Hazard Mapping System fire and smoke products [31], [32]. These official sources support a source-preserving Los Angeles County forecasting pipeline without requiring manually authored event records.

E. Explainable and Event-Aware Forecasting

Explainability is important in environmental forecasting because users need to understand why a model predicts a

rise or fall in pollution. SHAP values provide a model-agnostic framework for estimating feature contributions [33]. However, explanations based only on feature importance may miss temporal and contextual cues. Event-TimeRAF therefore generates explanations from structured evidence: recent PM2.5 persistence, weather flags, retrieved similar historical windows, retrieved event records, calendar context, and drift indicators. The explanation module uses templates grounded in stored evidence, not unrestricted language-model generation.

III. CRITICAL GAPS AND CORE CONTRIBUTIONS

A. Critical Gaps in Prior Studies

Existing forecasting approaches leave several gaps for event-sensitive urban air quality forecasting.

Gap 1: numerical history alone is not enough. Statistical, deep, and foundation forecasting models can learn temporal regularities, but PM2.5 changes often depend on weather, calendar behavior, and source-recorded environmental events. A model that ignores those factors can retrieve or learn patterns that are numerically similar but contextually different.

Gap 2: retrieval must be leakage-safe. Time-series retrieval can easily introduce look-ahead bias if retrieved candidates overlap the query input or target period. Environmental forecasting studies also risk using post-hoc event labels as if they were available before the forecast origin.

Gap 3: event availability is rarely audited. Public event databases are useful but not always operationally timestamped. NOAA Storm Events is official, but it does not provide a machine-readable publication timestamp. Without an explicit availability rule, event-aware forecasting claims may overstate real-time usability.

Gap 4: foundation-model gains need domain verification. TSFMs can be strong zero-shot forecasters, but their advantage is not guaranteed in a local PM2.5 task where supervised weather-calendar features are highly informative. Retrieval-augmented foundation-model claims therefore require a frozen-backbone baseline and a retrieval-augmented variant evaluated on the same test origins.

Gap 5: explanations should be evidence-grounded. Feature-importance scores alone do not explain whether a forecast is linked to weather shift, retrieved historical cases, event context, or drift. A deployable environmental forecasting workflow needs traceable evidence rather than unrestricted generated text.

B. Core Contributions

The proposed Event-TimeRAF framework addresses these gaps through the following components:

- *Context-aware retrieval knowledge base:* a source-preserving knowledge base that combines PM2.5 windows, weather summaries, calendar signatures, and NOAA Storm Events context.
- *Causal retrieval eligibility:* a strict candidate-filtering rule that allows a candidate only when its future target window finishes before the query lookback begins.
- *Audited event availability:* a data-readiness audit that records event source, delivery mode, source overlap,

TABLE I
MAIN SYMBOLS USED IN THE PROPOSED FORMULATION

Symbol	Meaning
L	Lookback window length
H	Forecast horizon
X_t	Historical PM2.5 sequence at time t
W_t	Weather covariates
C_t	Calendar/context features
R_t^{ts}	Retrieved time-series knowledge
R_t^{ev}	Retrieved event/news knowledge
D_t	Drift indicators
$\hat{Y}_{t+1:t+H}$	Predicted future PM2.5 values

category coverage, and whether the event timestamps are strict publication times or retrospective event-start assumptions.

- *Dual forecasting evaluation:* direct XGBoost feature-level fusion is evaluated alongside frozen Chronos-Bolt and validation-weighted retrieval-augmented Chronos fusion.
- *Evidence-based explanation layer:* forecast explanations are constructed from stored retrieval, weather, event, drift, and local model contribution evidence.

Unlike a full TimeRAF reproduction, this work does not train a dual-encoder retriever or implement internal Channel Prompting. Instead, it provides an audited domain adaptation of the retrieval-augmentation idea and tests whether the idea is useful under official-source PM2.5 forecasting constraints.

IV. PROBLEM FORMULATION

Let x_t denote the PM2.5 measurement at time t . For a forecast time t , the historical air-quality lookback window is

$$X_t = [x_{t-L+1}, x_{t-L+2}, \dots, x_t], \quad (1)$$

where L is the lookback length. Let W_t denote weather covariates aligned with the lookback window, C_t denote calendar features, and E_t denote event context retrieved from an event/news knowledge base. The target is the future PM2.5 sequence

$$Y_{t+1:t+H} = [x_{t+1}, x_{t+2}, \dots, x_{t+H}], \quad (2)$$

where H is the forecasting horizon.

The forecasting goal is to learn a function

$$\hat{Y}_{t+1:t+H} = F_\theta(X_t, W_t, C_t, R_t^{ts}, R_t^{ev}, D_t), \quad (3)$$

where R_t^{ts} is retrieved time-series knowledge, R_t^{ev} is retrieved event knowledge, and D_t contains concept-drift indicators. The implemented target setting is $L = 168$ hours and $H = 24$ hours.

The evaluation uses a time-based split:

$$\mathcal{T}_{train} < \mathcal{T}_{val} < \mathcal{T}_{test}, \quad (4)$$

where the earliest 70% of observations are used for training, the next 15% for validation, and the latest 15% for testing. Random splitting is not used because it would leak future temporal information into training.

Table I summarizes the essential notation.

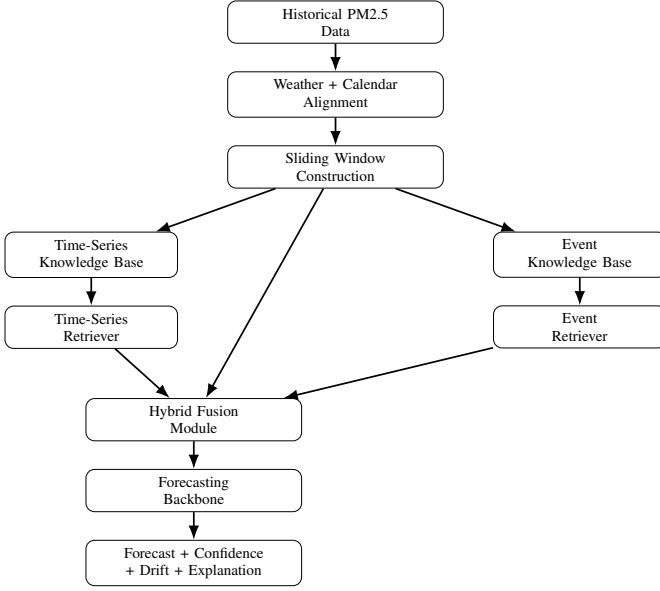


Fig. 1. Event-TimeRAF architecture used by the implemented pipeline.

V. METHODOLOGY

A. Overview

Fig. 1 shows the Event-TimeRAF pipeline. The first step is to align air-quality observations with weather, calendar, and availability-filtered event features. The system then creates sliding windows and a time-series knowledge base. The implemented hybrid retriever ranks historical windows using PM2.5, weather, calendar, and event-context similarity. A separate source-record lookup links recent event identifiers to explanations; it does not supply future event labels to the model. The retrieved evidence is combined with structured features to generate a forecast, uncertainty proxy, drift score, and explanation.

For each forecast origin t , Event-TimeRAF maps the aligned context into a forecast through

$$\hat{Y}_{t+1:t+H} = F_{\theta}(\phi(X_t, W_t, C_t, E_t), G_t, D_t), \quad (5)$$

where $\phi(\cdot)$ denotes the tabular feature construction, G_t denotes retrieved historical evidence, and D_t denotes drift evidence. The key methodological constraint is that every element of G_t and D_t must be available no later than the forecast origin.

B. Relationship to TimeRAF

TimeRAF retrieves top- k candidate time-series windows from a curated knowledge base and incorporates the retrieved candidates using Channel Prompting [9]. It uses a learnable dual-encoder retriever, dot-product similarity, a frozen TSFM backbone, and candidate augmentation during retriever training. Event-TimeRAF keeps the retrieval-augmentation principle, but applies it to Los Angeles County PM2.5 forecasting and expands the knowledge sources beyond numerical sequences.

The implemented system is not a complete clone of TimeRAF. Instead, it implements TimeRAF-style retrieval baselines using random retrieval, cosine similarity, and a

transparent hybrid similarity score. A dual-encoder MLP retriever, candidate augmentation, and full TSFM-based Channel Prompting are reserved for later extensions. The final run includes a frozen Chronos-Bolt gate and a validation-weighted forecast-level fusion between Chronos and retrieved historical forecasts.

C. Dataset Construction

Hourly PM2.5 observations are taken from US EPA Air-Data/AQS parameter 88101 for Los Angeles County [29]. The data audit selects one monitoring site using coverage and continuity criteria when a single site satisfies the coverage gate. In the final run, no single site is used as the target; the audited fallback is the hourly county median across official Los Angeles County AQS monitors. Timestamps are converted to *America/Los_Angeles* with explicit daylight-saving handling. Short input gaps may be filled only from past-available observations; targets are never interpolated. Weather covariates come from the nearest suitable NOAA Integrated Surface Database station [30]. Calendar features include hour, weekday, weekend, month, season, US federal holidays, and California holidays.

All normalization statistics are fitted on the training split only. The same normalization procedure is used for query windows and retrieved candidate windows, following the TimeRAF requirement that input and retrieved sequences use consistent preprocessing.

D. Time-Series Knowledge Base

The time-series knowledge base contains historical PM2.5 windows. Each record stores a window identifier, start time, end time, input window values, future window values, weather summary, calendar summary, normalized window vector, data split, and normalization identifier. Knowledge-base origins are separated by 192 hours, equal to the 168-hour input and 24-hour future combined. For every query, a candidate is eligible only when its complete future ends before the query input begins. Validation and test retrieval is additionally restricted to training-history windows. These rules exclude self-matches, near-duplicate windows, and overlap with the query context.

Formally, the knowledge base is

$$\mathcal{K} = \{(u_i, X_i, Y_i, m_i, c_i, e_i)\}_{i=1}^N, \quad (6)$$

where u_i is the candidate origin, X_i and Y_i are the candidate lookback and future windows, m_i is a weather summary, c_i is a calendar summary, and e_i is the event-context summary available at u_i . For a query origin t , the causally eligible set is

$$\mathcal{A}_t = \{i : u_i + H < t - L + 1, i \in \mathcal{T}_{train}\}, \quad (7)$$

for validation and test queries. The strict inequality in (7) ensures that a candidate's target period finishes before the query input period begins.

Given a query window X_t , the time-series retriever returns

$$R_t^{ts} = \{r_{t,1}^{ts}, r_{t,2}^{ts}, \dots, r_{t,k}^{ts}\}, \quad (8)$$

where k is the number of retrieved candidates. The default setting is $k = 8$, with sensitivity analysis for $k \in \{1, 4, 8, 16\}$.

E. Event Knowledge Base

The event knowledge base contains source-preserving records relevant to Los Angeles County. NOAA Storm Events is the required structured source, while NOAA Hazard Mapping System fire and smoke records remain an optional enrichment [31], [32]. Stored fields include event identifier, event start and end, information-availability assumption, location, category, source, source URL, title, and summary. Event categories include wildfire, smoke, excessive heat, high wind, heavy rain, flood, dust, and other weather-related records.

Model inputs represent events with hourly publication-assumption counts, rolling 24-hour, 72-hour, and 168-hour counts, category-specific 72-hour counts, active-event counts, and an event-burst ratio. Recent source event identifiers are retained only as explanation evidence. Every event used as a model input must be available by the forecast origin under a recorded availability rule. Evaluation distinguishes events published in the prior 72 hours, events active at the origin, and recorded event intervals overlapping the target horizon. Target overlap is used only as a post-hoc label and never as a model input. NOAA Storm Events does not provide a machine-readable publication timestamp; consequently, the implementation records event start as a retrospective availability assumption. Results that use this assumption are labeled sensitivity experiments rather than strict operational forecasts. Strict experiments require source records with genuine publication or issuance timestamps.

F. Hybrid Retrieval Score

The hybrid retrieval score combines time-series similarity, weather relevance, calendar relevance, and event relevance:

$$s(q, i) = \alpha s_{ts}(q, i) + \beta s_w(q, i) + \gamma s_c(q, i) + \delta s_e(q, i), \quad (9)$$

where s_{ts} is time-series similarity, s_w is weather similarity, s_c is calendar similarity, and s_e is event relevance. The implemented weights are $\alpha = 0.5$, $\beta = 0.2$, $\gamma = 0.1$, and $\delta = 0.2$.

The retrieved set is selected as

$$G_t = \text{TopK}_{i \in \mathcal{A}_t} s(t, i). \quad (10)$$

For retrieval-only baselines, the forecast is the uniform or similarity-weighted average of the retrieved candidates' future PM2.5 windows. For feature-level models, the retrieved futures and candidate metadata are summarized as mean, spread, similarity, and event-match features.

G. Fusion and Forecasting Backbone

The feature-level fusion method combines PM2.5 lag/rolling features, weather features, calendar features, retrieved time-series summary features, event features, and drift indicators. The tabular forecasting backbone uses 24 direct XGBoost regressors, one for each forecast horizon, because this design is efficient on Kaggle and provides a strong baseline. The foundation-model gate additionally evaluates frozen Chronos-Bolt and a validation-weighted forecast-level fusion between Chronos and retrieved historical forecasts. A later extension

may use an MLP residual fusion module inspired by Channel Prompting:

$$z_i = \text{Concat}(h_t, r_{t,i}), \quad (11)$$

$$h_t^* = h_t + \frac{1}{k} \sum_{i=1}^k \text{MLP}(z_i), \quad (12)$$

where h_t is the input representation and $r_{t,i}$ is the representation of the i -th retrieved candidate. This mirrors the TimeRAF idea of preserving the original input representation while adding uniformly averaged information extracted from retrieved candidates.

The frozen Chronos-Bolt gate uses forecast-level fusion rather than internal Channel Prompting. Let $\hat{Y}_t^C$ be the frozen Chronos forecast and $\hat{Y}_t^R$ be the retrieved historical forecast. The fused prediction is

$$\hat{Y}_t^{CR}(\omega) = \omega \hat{Y}_t^C + (1 - \omega) \hat{Y}_t^R, \quad (13)$$

where $\omega \in \{0, 0.25, 0.5, 0.75, 1.0\}$ is selected on the validation split. The final run selected $\omega = 0.75$.

H. Concept Drift Detection

The implementation represents operational distribution shift using transparent indicators rather than claiming perfect concept-drift identification. The indicators are recent mean shift, recent variance shift, retrieval similarity drop, weather-pattern shift, and event-burst ratio. A simple drift score can be defined as

$$D_t = \lambda_1 z_{\mu,t} + \lambda_2 z_{\sigma,t} + \lambda_3 (1 - \bar{s}_{ts,t}) + \lambda_4 z_{w,t} + \lambda_5 b_{event,t}, \quad (14)$$

where $z_{\mu,t}$ and $z_{\sigma,t}$ are normalized mean and variance shifts, $\bar{s}_{ts,t}$ is average retrieval similarity, $z_{w,t}$ is weather-pattern shift, and $b_{event,t}$ is an event-burst score. The implemented components are robustly scaled against training data and averaged with equal weight; the alert threshold is selected from validation scores.

I. Explanation Generation

The explanation module generates evidence-based explanations. It uses signed local XGBoost contributions averaged across the 24 direct horizon models, weather flags, retrieved historical windows, retrieved event records, calendar context, forecast direction, component-level drift evidence, and an uncertainty proxy combining retrieved-trajectory spread with validation residual error. It does not generate explanations from an unrestricted language model. Each explanation is traceable to stored machine-readable evidence.

Table II summarizes the methodological relation between TimeRAF and Event-TimeRAF.

VI. EXPERIMENTAL SETUP

This section reports the executed final Kaggle experiment. The run identifier is 20260723T112033170131Z. All model-selection choices were frozen before reading the final test results, and the results notebook rejected non-final or mixed-run artifacts.

TABLE II
METHODOLOGICAL COMPARISON BETWEEN TIMERAF AND
EVENT-TIMERAF

Aspect	TimeRAF	Event-TimeRAF
Target domain	General zero-shot TSF	Los Angeles County PM2.5 forecasting
Knowledge base	Time-series windows	PM2.5 windows, weather, calendar, events
Retriever	Learnable dual encoder	Random, cosine, and hybrid similarity
Integration	Channel Prompting	Feature fusion; forecast-level TSFM retrieval fusion
Backbone	Frozen TSFM	Direct XGBoost; frozen Chronos-Bolt gate
Explanations	Retrieved-case visualization	Evidence-based forecast explanation
Drift handling	Not central	Explicit drift indicators

A. Data Sources and Audit

The study period is 2019–2024. PM2.5 observations were taken from US EPA AirData/AQS parameter 88101 for Los Angeles County [29]. No single monitoring site was used as the final target; the audited target is the hourly county median across official Los Angeles County AQS monitors. This fall-back selected 06–037–COUNTY, with 52,602 observed target hours and observed PM2.5 coverage of 0.9997. The PM2.5 unit was consistently recorded as micrograms per cubic meter.

Weather covariates were taken from NOAA Global Hourly/ISD through the official NOAA Open Data Dissemination public buckets [30]. The selected station was 72297023129, Long Beach / Daugherty Field / Airport, 10.28 km from the PM2.5 target centroid. The complete weather-feature coverage after allowed past-only short-gap filling was 0.9942, and the raw measured core coverage was 0.9905.

Environmental events were loaded from hash-verified NOAA Storm Events annual detail archives [31]. Because Kaggle could not reliably download NCEI Storm Events directly, the experiment used a private delivery cache containing unchanged NOAA annual archives and a source manifest with official URLs, byte counts, and SHA-256 hashes. The audit retained the delivery mode `verified_official_noaa_cache`. The event audit found 117 event days, 2,192 source-overlap days, and three qualifying event categories. NOAA Storm Events does not expose machine-readable publication timestamps; therefore, event-aware outputs use the recorded event start time as a retrospective availability assumption. They are not strict real-time event forecasts.

B. Data and Code Availability

The raw public records are available from EPA AirData, NOAA Global Hourly/ISD, and NOAA Storm Events [29]–[31]. The experiment used a private Kaggle delivery cache for NOAA Storm Events only because direct NCEI access was unavailable in the notebook environment; every cached NOAA

TABLE III
FINAL DATA-READINESS AUDIT

Gate	Value	Status
PM2.5 observed coverage	0.9997	Pass
Weather complete coverage	0.9942	Pass
Event days	117	Pass
Event overlap days	2,192	Pass
Qualifying event categories	3	Pass
Strict event availability	False	Caveat

file was verified against the generated source manifest. Source code is not publicly available.

C. Forecasting Task and Split

The forecasting task predicts PM2.5 from one to 24 hours ahead using a 168-hour lookback window. The split is chronological: the earliest 70% of valid windows are used for training, the next 15% for validation, and the latest 15% for testing. The final test set contains 7,199 forecast origins and 172,776 forecasted hourly target points. Random splitting is not used.

D. Models

The evaluated models are persistence, daily seasonal naive, weekly seasonal naive, a direct XGBoost model using PM2.5 features, a direct XGBoost context model using PM2.5, weather, and calendar features, random retrieval, cosine retrieval, XGBoost with cosine retrieval features, Event-TimeRAF without drift features, full Event-TimeRAF with event and drift features, frozen Chronos-Bolt, and retrieval-augmented Chronos-Bolt forecast fusion. The Chronos checkpoint is `amazon/chronos-bolt-small`; it completed successfully in the final run, so the foundation-model publication gate is satisfied.

E. Implementation and Hyperparameters

The final run was executed on Kaggle using Python 3.12.13 on Linux 6.12.90. The runtime was 1,576.87 seconds. Key package versions recorded in the run manifest include NumPy 2.0.2, pandas 2.3.3, scikit-learn 1.6.1, XGBoost 3.2.0, PyTorch 2.10.0+cu128, and chronos-forecasting 2.3.1. Table VI reports the principal hyperparameters.

F. Metrics and Statistical Checks

The primary metrics are MSE, MAE, RMSE, MAPE, sMAPE, and R^2 . MSE is retained for comparability with the TimeRAF base paper, while MAE and RMSE provide more interpretable PM2.5 error scales. Subset metrics are reported only when at least 50 forecast origins are available. The event subset contains 532 origins, the active-event subset contains 293 origins, and the drift subset contains 560 origins, so all pre-specified subsets pass the minimum-size gate. Pairwise comparisons use paired block bootstrap intervals with 24-hour blocks and 500 resamples.

TABLE IV
DATASET AND SOURCE SUMMARY

Data group	Source	Final scope	Use in pipeline
PM2.5	EPA AirData/AQS hourly parameter 88101	Los Angeles County, 2019–2024; 52,602 observed target hours	Target sequence and lag/rolling features
Weather	NOAA Global Hourly/ISD, station 72297023129	2019–2024; 0.9942 usable core coverage	Temperature, humidity, wind, pressure, and rain features
Events	NOAA Storm Events annual detail archives	204 Los Angeles records; 117 event days	Retrospective event-start features and explanation evidence
Calendar	Generated from timestamps and public holiday rules	2019–2024 hourly calendar	Hour, weekday, weekend, month, season, and holiday features

TABLE V
EVALUATED MODEL FAMILIES

ID	Description
M00–M02	Persistence and seasonal naive baselines
M03–M04	XGBoost PM2.5-only and context models
M05–M06	Random and cosine retrieval-only baselines
M07	XGBoost with cosine retrieval features
M08–M09	Event-TimeRAF without/with drift indicators
M10–M11	Frozen Chronos-Bolt and retrieval fusion

G. Position Relative to Existing Studies

Direct numerical comparison with existing TSFM and air-quality papers is not valid because they use different datasets, forecast horizons, variables, and evaluation protocols. Table VII therefore compares methodological coverage rather than claiming cross-paper performance dominance.

VII. RESULTS AND DISCUSSION

Table VIII reports the final test-set results. The strongest overall model is the XGBoost context baseline, which uses PM2.5 history, weather, and calendar features. It obtains MSE 26.185 and MAE 3.125. The full Event-TimeRAF XGBoost variant obtains MSE 26.712 and MAE 3.149. Therefore, this run does not support a claim that the full Event-TimeRAF XGBoost variant outperforms the strongest non-retrieval context baseline on the all-origin test set. It does, however, substantially outperform the persistence and seasonal naive baselines.

The frozen Chronos-Bolt publication gate completed successfully. The validation fusion search selected a retrieval-fusion weight of 0.75. On the test set, retrieval-augmented Chronos reduces MSE from 28.941 to 28.709, an improvement of about 0.8% relative to frozen Chronos-Bolt. The paired block-bootstrap comparison gives an MSE difference of -0.233 with a 95% interval from -0.665 to 0.152, so the direction is favorable but not conclusive under this uncertainty estimate. MAE changes from 3.205 to 3.209, which is practically flat.

Table IX summarizes the main ablation comparisons. The clearest supported gain is from adding weather and calendar context to the PM2.5-only XGBoost model: MSE decreases by 3.764 and the confidence interval is fully below zero. Cosine retrieval improves over random retrieval in the retrieval-only setting. In contrast, adding cosine retrieval features to the already strong XGBoost context model does not improve




Fig. 2. MSE by forecast horizon for the final test run.

the all-origin result in this run. Event-hybrid retrieval and drift indicators also do not show a clear overall improvement beyond the context baseline.

The k sensitivity results in Table X show that the retrieval-only cosine baseline improves as more neighbors are used within the tested range. The best retrieval-only configuration among $k \in \{1, 4, 8, 16\}$ is $k = 16$, with MSE 36.488. The final Event-TimeRAF feature models used the pre-specified default $k = 8$.

Table XI reports representative event and drift subsets for the strongest context model, the full Event-TimeRAF model, and the two Chronos variants. All subsets pass the minimum-size gate. The event subset has lower MSE than the non-event subset for these models, which should not be interpreted as events being easier in general; it reflects the particular 2024 test-period event hours retained by the audit. On drift-flagged origins, the full Event-TimeRAF model is worse than the context baseline, while retrieval-augmented Chronos is slightly better than frozen Chronos by MSE. This again indicates that the present event/drift feature fusion is not sufficient to dominate the context baseline.

Fig. 2 and Fig. 3 show horizon-wise test errors. Fig. 4 summarizes retrieval behavior, and Fig. 5 shows the drift-score distribution across the test period. These plots are generated directly from the saved final-run artifacts and share run identifier 20260723T112033170131Z.

A. Explainability Analysis

Explainability is implemented as a source-grounded evidence record rather than as a post-hoc free-text interpretation. For each evaluated forecast origin, the pipeline stores the run

TABLE VI
FINAL RUN HYPERPARAMETERS AND RUNTIME SETTINGS

Component	Setting
Study period	2019–2024
Lookback and horizon	$L = 168$ hours, $H = 24$ hours
Split	70% train, 15% validation, 15% test, chronological
Retrieval	$k = 8$, sensitivity for $k \in \{1, 4, 8, 16\}$
Knowledge-base stride	192 hours
Hybrid retrieval weights	time-series 0.5, weather 0.2, calendar 0.1, event 0.2
XGBoost	350 estimators, max depth 6, learning rate 0.05, subsample 0.85, colsample 0.85, $\lambda = 1.0$
Drift detector	24-hour recent window, 168-hour reference window, 0.90 validation quantile threshold
Chronos checkpoint	amazon/chronos-bolt-small, batch size 64
Chronos fusion weights tested	$\{0, 0.25, 0.5, 0.75, 1.0\}$; selected 0.75
Bootstrap	500 paired block-bootstrap resamples, 24-hour blocks

TABLE VII
METHODOLOGICAL POSITIONING AGAINST RELATED FORECASTING WORK

Method family	TSFM	Retrieval	Events	Explainability	Comment
PatchTST-style Transformers [21]	No	No	No	Limited	Strong sequence modeling, but event context is not central.
Chronos [5]	Yes	No	No	Limited	Frozen zero-shot baseline used in this study.
TimeRAF [9]	Yes	Yes	No	Retrieval cases	Base retrieval-augmented TSFM method; does not target official event-aware PM2.5.
XGBoost context baseline [10]	No	No	No	Feature importance	Strong supervised tabular baseline for local PM2.5.
Event-TimeRAF	Yes	Yes	Yes	Evidence-grounded	Adds official event audit, causal retrieval, drift indicators, and Chronos retrieval fusion.

TABLE VIII
FINAL OVERALL TEST RESULTS FOR 24-HOUR PM2.5 FORECASTING

Model	MSE	MAE	RMSE	R^2
M00 persistence	47.649	4.137	6.903	-0.130
M01 daily seasonal	48.706	4.057	6.979	-0.155
M02 weekly seasonal	65.606	5.123	8.100	-0.556
M03 XGBoost PM2.5	29.949	3.368	5.473	0.290
M04 XGBoost context	26.185	3.125	5.117	0.379
M05 random retrieval	42.367	4.257	6.509	-0.005
M06 cosine retrieval	38.083	3.940	6.171	0.097
M07 XGBoost cosine	26.672	3.151	5.164	0.367
M08 Event-TimeRAF no drift	26.677	3.145	5.165	0.367
M09 Event-TimeRAF full	26.712	3.149	5.168	0.367
M10 frozen Chronos-Bolt	28.941	3.205	5.380	0.314
M11 Chronos + retrieval	28.709	3.209	5.358	0.319

TABLE IX
SELECTED PAIRED BLOCK-BOOTSTRAP ABLATION RESULTS

Comparison	MSE diff.	95% CI low	95% CI high
M04 minus M03, weather/calendar	-3.764	-5.856	-2.071
M06 minus M05, cosine/random retrieval	-4.284	-5.862	-2.759
M07 minus M04, retrieval features	0.487	0.015	1.028
M08 minus M07, event hybrid	0.005	-0.292	0.275
M09 minus M08, drift indicators	0.035	-0.101	0.161
M09 minus M04, full Event-TimeRAF	0.528	-0.026	1.182
M11 minus M10, Chronos retrieval	-0.233	-0.665	0.152

identifier, window identifier, origin timestamp, retrieved historical evidence identifiers, event evidence identifiers, leading feature effects, drift components, retrieval spread, validation residual MAE, uncertainty proxy, drift score, drift flag, and event-availability mode. The final run produced 7,199 explana-

tion records in `explanations.parquet`. A representative record identifies the strongest local contribution fields as the current PM2.5 value and recent rolling PM2.5 means, links the prediction to three retrieved historical windows, links recent source-recorded storm events when available, and marks

TABLE X
COSINE RETRIEVAL SENSITIVITY TO k

k	MSE	MAE	RMSE
1	65.129	5.133	8.070
4	41.237	4.134	6.422
8	38.083	3.940	6.171
16	36.488	3.851	6.040




Fig. 3. MAE by forecast horizon for the final test run.




Fig. 4. Retrieval diagnostics for the final test run.




Fig. 5. Drift-score diagnostics for the final test run.




Fig. 6. Representative forecast case generated from the final test run.

distribution shift when the drift detector is active.

Table XII summarizes the explanation layer. This design supports auditability because every explanatory statement can be traced back to a model feature, a retrieved time-series

window, an official event record, or a drift component. It does not prove causality: event evidence is reported as contextual support, and the drift flag is treated as an uncertainty warning.

B. Computational Cost and Reproducibility

The final publication-mode run completed in 1,576.9 seconds in a Kaggle Linux notebook environment using Python 3.12.13. The recorded package versions include NumPy 2.0.2, pandas 2.3.3, scikit-learn 1.6.1, XGBoost 3.2.0, PyTorch 2.10.0+cu128, and chronos-forecasting 2.3.1. The run manifest records the configuration SHA-256 hash and artifact checksums. The CPU, GPU, and RAM model identifiers were not saved in the manifest; therefore this paper reports runtime and software versions but does not claim hardware-normalized efficiency.

The main interpretation is that conventional supervised context features are very strong for this Los Angeles County 24-hour PM_{2.5} task. Retrieval is useful as a standalone historical-pattern baseline and gives a small favorable MSE change when fused with frozen Chronos-Bolt, but the present feature-level Event-TimeRAF design does not beat the best XGBoost context baseline overall. This negative result is important for the final paper: Event-TimeRAF should be presented as an implemented and audited event-aware retrieval framework with mixed empirical effects, not as a uniformly superior forecasting model.

Finally, all event-aware interpretations must be read with the availability caveat. NOAA Storm Events records are official and hash-verified, but the source does not provide machine-readable publication timestamps. The experiment records event start as the availability time, so event-related results are retrospective sensitivity results. A strict operational study would require event or smoke sources with genuine issue or publication times, such as an audited NOAA HMS cache or alert product archive.

VIII. LIMITATIONS AND FUTURE WORK

The most important limitation is event availability. NOAA Storm Events records are official and source-verified, but the public detail files do not provide machine-readable publication timestamps. The event-aware component therefore uses event start time as the availability proxy. This makes the reported event-aware results retrospective sensitivity results rather than strict operational nowcasts. A production deployment would require sources with auditable issue times, such as alert archives, smoke products, or an HMS cache with verified timestamps.

The second limitation is external validation. The final experiment evaluates one carefully audited setting: Los Angeles County PM_{2.5} from 2019–2024. This is appropriate for a scoped Kaggle implementation, but it does not establish geographic generalization. The faculty-quality target asks for validation on a second dataset; for this study, that remains future work rather than a result. A stronger version should repeat the same protocol on another U.S. county or city with different meteorology, emissions, and event profiles.

TABLE XI
REPRESENTATIVE EVENT AND DRIFT SUBSET RESULTS

Subset	Model	Origins	MSE	MAE
Non-event	M04 XGBoost context	6,667	27.520	3.179
Event	M04 XGBoost context	532	9.451	2.450
Non-event	M09 Event-TimeRAF full	6,667	27.942	3.187
Event	M09 Event-TimeRAF full	532	11.301	2.666
Non-drift	M04 XGBoost context	6,639	26.107	3.088
Drift	M04 XGBoost context	560	27.110	3.562
Non-drift	M09 Event-TimeRAF full	6,639	26.330	3.088
Drift	M09 Event-TimeRAF full	560	31.244	3.872
Drift	M10 frozen Chronos-Bolt	560	28.333	3.476
Drift	M11 Chronos + retrieval	560	28.198	3.580

TABLE XII
EXPLANATION EVIDENCE STORED BY EVENT-TIMERAF

Evidence field	Interpretation role
Top feature effects	Local model drivers from the direct fore-caster.
Retrieved windows	Historical analogues used by the retrieval module.
Event evidence IDs	Official storm-event context linked to the origin.
Drift components	Mean, variance, weather, similarity, and event-burst shift signals.
Uncertainty proxy	Combined retrieval spread and validation residual error.

TABLE XIII
FINAL-RUN COMPUTATIONAL AND REPRODUCIBILITY SUMMARY

Item	Recorded value
Run identifier	20260723T112033170131Z
Runtime	1,576.9 seconds
Platform	Kaggle Linux notebook
Python	3.12.13
Forecast horizon	24 hours
Lookback window	168 hours
Bootstrap resamples	500
Chronos checkpoint	amazon/chronos-bolt-small
Code availability	Source code is not publicly available.

The third limitation concerns model integration. Event-TimeRAF currently uses retrieved trajectories and event context as feature-level evidence, with a frozen Chronos-Bolt gate for foundation-model validation. This is weaker than a learnable retrieval-and-prompting architecture that directly conditions the foundation model. Future work should test a learnable retriever, richer event-type embeddings, calibrated uncertainty estimation, and a closer adaptation of TimeRAF-style Channel Prompting for environmental forecasting.

Finally, the present manuscript uses verified references and official data sources only. The reference list should be expanded with additional recent 2023–2026 air-quality forecasting, time-series foundation model, retrieval, and explainability studies before journal submission.

IX. CONCLUSION

This paper presented Event-TimeRAF, an event-aware retrieval-augmented framework for explainable Los Angeles

County PM2.5 forecasting under distribution shift. The implemented pipeline combines official EPA PM2.5 observations, NOAA weather, hash-verified NOAA Storm Events records, leakage-safe time-series retrieval, direct XGBoost forecasters, drift indicators, evidence-based explanations, and a frozen Chronos-Bolt publication gate. In the final 2019–2024 Kaggle experiment, the strongest all-origin model is the XGBoost weather-calendar context baseline with MSE 26.185, while the full Event-TimeRAF XGBoost variant obtains MSE 26.712. Retrieval improves the standalone cosine baseline over random retrieval and gives a small favorable MSE change when fused with frozen Chronos-Bolt, but the full feature-level design does not outperform the strongest supervised context baseline. The study therefore positions Event-TimeRAF as an auditable event-aware retrieval framework with mixed empirical effects, and it identifies stricter event availability, external validation, and deeper foundation-model integration as the main directions for the next version.

REFERENCES

- [1] World Health Organization, “WHO global air quality guidelines: Particulate matter (PM2.5 and PM10), ozone, nitrogen dioxide, sulfur dioxide and carbon monoxide,” 2021. [Online]. Available: <https://www.who.int/publications/i/item/9789240034228>
- [2] A. Das, W. Kong, R. Sen, and Y. Zhou, “A decoder-only foundation model for time-series forecasting,” in *Proceedings of the 41st International Conference on Machine Learning*, 2024.
- [3] G. Woo, C. Liu, A. Kumar, C. Xiong, S. Savarese, and D. Sahoo, “Unified training of universal time series forecasting transformers,” 2024.
- [4] M. Goswami, K. Szafer, A. Choudhry, Y. Cai, S. Li, and A. Dubrawski, “MOMENT: A family of open time-series foundation models,” 2024.
- [5] A. F. Ansari, L. Stella, C. Turkmen, X. Zhang, P. Mercado, H. Shen, O. Shchur, S. S. Rangapuram, S. P. Arango, S. Kapoor, J. Zschiegner, D. C. Maddix, H. Wang, M. W. Mahoney, K. Torkkola, A. G. Wilson, M. Bohlke-Schneider, and Y. Wang, “Chronos: Learning the language of time series,” 2024.
- [6] V. Ekambaram, A. Jati, P. Dayama, S. Mukherjee, N. H. Nguyen, W. M. Gifford, C. Reddy, and J. Kalagnanam, “Tiny time mixers (TTMs): Fast pre-trained models for enhanced zero/few-shot forecasting of multivariate time series,” 2024.
- [7] P. Lewis, E. Perez, A. Piktus, F. Petroni, V. Karpukhin, N. Goyal, H. Kuttler, M. Lewis, W. tau Yih, T. Rocktaschel, S. Riedel, and D. Kiela, “Retrieval-augmented generation for knowledge-intensive NLP tasks,” in *Advances in Neural Information Processing Systems*, vol. 33, 2020, pp. 9459–9474.
- [8] V. Karpukhin, B. Oguz, S. Min, P. Lewis, L. Wu, S. Edunov, D. Chen, and W. tau Yih, “Dense passage retrieval for open-domain question answering,” in *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing*, 2020, pp. 6769–6781.

- [9] H. Zhang, C. Xu, Y.-F. Zhang, Z. Zhang, L. Wang, and J. Bian, "TimeRAF: Retrieval-augmented foundation model for zero-shot time series forecasting," *IEEE Transactions on Knowledge and Data Engineering*, vol. 37, no. 9, pp. 5654–5665, 2025.
- [10] T. Chen and C. Guestrin, "XGBoost: A scalable tree boosting system," in *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 2016, pp. 785–794.
- [11] G. Ke, Q. Meng, T. Finley, T. Wang, W. Chen, W. Ma, Q. Ye, and T.-Y. Liu, "LightGBM: A highly efficient gradient boosting decision tree," in *Advances in Neural Information Processing Systems*, vol. 30, 2017.
- [12] S. Hochreiter and J. Schmidhuber, "Long short-term memory," *Neural Computation*, vol. 9, no. 8, pp. 1735–1780, 1997.
- [13] K. Cho, B. van Merriënboer, C. Gulcehre, D. Bahdanau, F. Bougares, H. Schwenk, and Y. Bengio, "Learning phrase representations using RNN encoder–decoder for statistical machine translation," in *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing*, 2014, pp. 1724–1734.
- [14] H. Zhou, S. Zhang, J. Peng, S. Zhang, J. Li, H. Xiong, and W. Zhang, "Informer: Beyond efficient transformer for long sequence time-series forecasting," in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 35, no. 12, 2021, pp. 11 106–11 115.
- [15] H. Wu, J. Xu, J. Wang, and M. Long, "Autoformer: Decomposition transformers with auto-correlation for long-term series forecasting," in *Advances in Neural Information Processing Systems*, vol. 34, 2021.
- [16] T. Zhou, Z. Ma, Q. Wen, X. Wang, L. Sun, and R. Jin, "FEDformer: Frequency enhanced decomposed transformer for long-term series forecasting," in *Proceedings of the 39th International Conference on Machine Learning*, ser. Proceedings of Machine Learning Research, vol. 162, 2022, pp. 27 268–27 286.
- [17] Y. Zhang and J. Yan, "Crossformer: Transformer utilizing cross-dimension dependency for multivariate time series forecasting," in *Proceedings of the 11th International Conference on Learning Representations*, 2023.
- [18] Y. Liu, T. Hu, H. Zhang, H. Wu, S. Wang, L. Ma, and M. Long, "iTransformer: Inverted transformers are effective for time series forecasting," in *Proceedings of the 12th International Conference on Learning Representations*, 2024.
- [19] H. Wu, T. Hu, Y. Liu, H. Zhou, J. Wang, and M. Long, "TimesNet: Temporal 2d-variation modeling for general time series analysis," in *Proceedings of the 11th International Conference on Learning Representations*, 2023.
- [20] A. Zeng, M. Chen, L. Zhang, and Q. Xu, "Are transformers effective for time series forecasting?" in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 37, no. 9, 2023, pp. 11 121–11 128.
- [21] Y. Nie, N. H. Nguyen, P. Sinthong, and J. Kalagnanam, "A time series is worth 64 words: Long-term forecasting with transformers," in *Proceedings of the 11th International Conference on Learning Representations*, 2023.
- [22] Y. Liang, H. Wen, Y. Nie, Y. Jiang, M. Jin, D. Song, S. Pan, and Q. Wen, "Foundation models for time series analysis: A tutorial and survey," 2024.
- [23] J. Ye, W. Zhang, K. Yi, Y. Yu, Z. Li, J. Li, and F. Tsung, "A survey of time series foundation models: Generalizing time series representation with large language models," 2024.
- [24] K. Rasul, A. Ashok, A. R. Williams, H. Ghonia, R. Bhagwatkar, A. Khorasani, M. J. D. Bayazi, G. Adamopoulos, R. Riachi, N. Hassen, M. Biloš, S. Garg, A. Schneider, N. Chapados, A. Drouin, V. Zantedeschi, Y. Nevmyvaka, and I. Rish, "Lag-llama: Towards foundation models for probabilistic time series forecasting," 2024.
- [25] S. Zhou, W. Wang, L. Zhu, Q. Qiao, and Y. Kang, "Deep-learning architecture for PM2.5 concentration prediction: A review," *Environmental Science and Ecotechnology*, vol. 21, p. 100400, 2024.
- [26] Z. Zhang and S. Zhang, "Modeling air quality PM2.5 forecasting using deep sparse attention-based transformer networks," *International Journal of Environmental Science and Technology*, vol. 20, pp. 13 535–13 550, 2023.
- [27] M. Pandey, V. Jain, N. Godhani, S. N. Tripathi, and P. Rai, "Spatio-temporal forecasting of PM2.5 via spatial-diffusion guided encoder-decoder architecture," 2024.
- [28] J. Zuo, W. Li, M. Baldo, and H. Hacid, "Opportunistic air quality monitoring and forecasting with expandable graph neural networks," 2023.
- [29] US Environmental Protection Agency, "AirData: Pre-generated air quality data files," https://aqs.epa.gov/aqsweb/airdata/download_files.html, 2026, accessed: 2026-07-17.
- [30] NOAA National Centers for Environmental Information, "Global hourly—integrated surface database," <https://www.ncei.noaa.gov/products/land-based-station/integrated-surface-database>, 2026, accessed: 2026-07-17.
- [31] —, "Storm events database: Bulk data download," <https://www.ncei.noaa.gov/stormevents/ftp.jsp>, 2026, accessed: 2026-07-17.
- [32] NOAA Office of Satellite and Product Operations, "Hazard mapping system fire and smoke product," <https://ospo.noaa.gov/products/land/hms.html>, 2026, accessed: 2026-07-17.
- [33] S. M. Lundberg and S.-I. Lee, "A unified approach to interpreting model predictions," in *Advances in Neural Information Processing Systems*, vol. 30, 2017.